



Learning with Covariance Matrices: Foundations and Applications to Network Neuroscience

Saurabh Sihag¹, Gonzalo Mateos², Elvin Isufi³, and Alejandro Ribeiro⁴

¹University at Albany – ssihag@albany.edu

²University of Rochester – gmateosb@ece.rochester.edu

³Delft University of Technology – e.isufi-1@tudelft.nl


⁴University of Pennsylvania – aribeiro@seas.upenn.edu

European Signal Processing Conference
(EUSIPCO), 2025




Covariance Matrix


- Covariance matrix captures the **redundancies** between data points (features)
 - **Brain datasets:** some areas of the brain activate together
 - **Financial datasets:** stock prices fluctuate in tandem
 - **Traffic datasets:** traffic volume is correlated across intersections



Brain



Finance



Traffic

Covariance Matrix

➤ Evaluating a covariance matrix

- Consider a random variable $\mathbf{x} \in \mathbb{R}^m$
- The covariance is

$$\mathbf{C} = \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^\top], \text{ where } \boldsymbol{\mu} = \mathbb{E}[\mathbf{x}]$$

- In practice, we have **sample** covariance matrix (an estimate)

$$\hat{\mathbf{C}} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^\top, \text{ where } \hat{\boldsymbol{\mu}} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$$

n: number of samples (size of a dataset)




Covariance Matrix

- Covariance matrix encodes **redundancies** between different features in data


Covariance matrix
(2-feature dataset)

$\sigma^2(r_1)$	$\sigma(r_1, r_2)$
$\sigma(r_1, r_2)$	$\sigma^2(r_2)$

Low redundancy
(smaller $\sigma(r_1, r_2)$)



High redundancy
(higher $\sigma(r_1, r_2)$)



$\sigma(r_1, r_2)$ = how features r_1 and r_2 vary with respect to each other


Covariance matrices are widespread in signal processing and machine learning

➤ Principal component analysis (PCA)

- Eigenvectors of the covariance matrix form principal components (PCs)
- PCs inform the shape of a dataset (directions of variance)

Given sample \mathbf{x} and eigendecomposition $\hat{\mathbf{C}} = \hat{\mathbf{V}}\hat{\Lambda}\hat{\mathbf{V}}^T$,

$$\text{PCA transform: } \tilde{\mathbf{x}} = \hat{\mathbf{V}}^T \mathbf{x}$$




PCA transform in ML

- Unsupervised learning (dim. reduction)
- Supervised learning (regression, classification)

Covariance matrices are widespread in signal processing and machine learning


➤ Covariance matrices are leveraged as **graphical** representations of data

- A graph $G = (V, E, W)$
 - Set of nodes V
 - Set of edges E
 - A weight function W
- Covariance matrix is a **fully connected graph**,
 - nodes are the features
 - edges associated with pairwise covariance values




Covariance matrices are widespread in signal processing and machine learning


- Covariance matrices as **graphical** representations; used in graph neural nets



Brain connectome [Li, et al. 2021]




Socio-economic networks [Leo, et al. 2016]



Learning with covariance matrices: Challenges

- Sample covariance matrix is estimate from **finite** data
- ML model is trained on **training** dataset, deployed on **test** dataset
- Statistical spaces defined by **training** and **test** data may not align perfectly




Learning with covariance matrices: Challenges

- Sample covariance matrix is estimate from **finite** data
- ML model is trained on **training** dataset, deployed on **test** dataset
- Statistical spaces defined by **training** and **test** data may not align perfectly

Challenge 1 (stability)



Are inference outcomes **stable** to perturbations in covariance matrix (finite sample effect)?



Learning with covariance matrices: Challenges



- Datasets capture information about same phenomenon at **different scales**

Dataset with m_1 features



Covariance matrix \mathbf{C}_{m_1}
(size $m_1 \times m_1$)


Dataset with m_2 features





Covariance matrix \mathbf{C}_{m_2}
(size $m_2 \times m_2$)

Learning with covariance matrices: Challenges


- Datasets capture information about same phenomenon at **different scales**



Covariance matrix \mathbf{C}_{m_1}
(size $m_1 \times m_1$)



Covariance matrix \mathbf{C}_{m_2}
(size $m_2 \times m_2$)



Challenge 2 (transferability)

Can the redundancy in covariance matrices of datasets of different sizes be exploited?

Learning with covariance matrices: A GSP approach

- Signal and information processing is about exploiting signal structure
- Graph signal processing (GSP): broaden classical signal processing to graphs



Graph Signal Processing: Overview, Challenges, and Applications

This article presents methods to process data associated to graphs (graph signals) extending techniques (transforms, sampling, and others) that are used for conventional signals.

By ANTONIO ORTEGA¹, Fellow IEEE, PASCAL FROSSARD, Fellow IEEE, JELENA KOVACIĆ, Fellow IEEE, JOSE M. F. MOURA², Fellow IEEE, and PIERRE VANDERHEYDEN

ABSTRACT | Research in graph signal processing (GSP) aims to develop tools for processing data defined on irregular graph domains. In this paper, we first provide an overview of core ideas in GSP and their connection to conventional digital signal processing, along with a brief historical perspective to highlight how concepts recently developed in GSP build on top of prior research in other areas. We then summarize recent advances in developing basic GSP tools, including methods for sampling, filtering, or graph learning. Next, we review progress in several application areas using GSP, including processing and analysis of sensor network data, biological data, and applications to image processing and machine learning.

KEYWORDS | Graph signal processing (GSP); network science and graphs; sampling; signal processing

I. INTRODUCTION AND MOTIVATION

Data is all around us, and massive amounts of it. Almost every aspect of human life is now being recorded at all levels: from the marking and recording of processes inside the cells starting with the advent of fluorescent markers, to our personal data through health monitoring devices and apps, financial and banking data, our social networks, mobility and traffic patterns, marketing preferences, fads, and many more. The complexity of such networks [1] and interactions mean that the data now reside on irregular and complex structures that do not lend themselves to standard tools.

The term *graph signal processing* was coined a decade ago in the technical works [1]–[4]. Since then, many papers have been published, GSP-related problems have drawn significant attention, not only within the SP community [5] but also in machine learning (ML) venues, where research in graph-based learning has increased significantly [6]. Graph signals are well-suited to model measurements/information/data associated with (indexed by) a set where 1) the elements of the set belong to the same class (regions of the cerebral cortex, members of a social network, weather stations across a continent); 2) there exists a relation (physical or functional) of proximity, influence, or association among the different elements of that set; and 3) the strength of such a relation among the pairs of elements is not homogeneous. In some scenarios, the supporting graph is a physical, technological, social, information, or biological network where the links can be explicitly observed. In many other cases, the graph is implicit, capturing some notion of dependence or similarity across nodes, and the links must be inferred from the data themselves. As a result, GSP is a broad framework that encompasses and extends classical SP methods, tools, and algorithms to application domains of the modern technological world, including social, transportation, communication,

Graphs offer the ability to model such data and complex interactions among them. For example, users on Twitter can be modeled as nodes while their friend connections can be modeled as edges. This paper explores adding attributes to such nodes and modeling those as signals on a graph; for example, year of graduation in a social network, temperature in a given city on a given day in a weather network, etc. Doing so requires us to extend classical signal processing concepts and tools such as Fourier transform, filtering, and frequency response to data residing on graphs. It also leads us to tackle complex tasks such as sampling in a principled way. The field that gathers all these questions under a common umbrella is graph signal processing (GSP) [2], [3].

While the field of GSP is still in its infancy, its impact will be large in the future. Let us assume for now that a graph signal is a set of values residing on a set of nodes. These nodes are connected via (possibly weighted) edges. As in classical signal processing, such signals can stem from a variety of domains; unlike in classical signal processing, however, the underlying graphs can tell a fair amount about those signals through their structure. Different types of graphs model different types of networks that these nodes represent.

Typical graphs that are used to represent common real-world data include Erdős-Rényi graphs, ring graphs, random geometric graphs, small-world graphs, power-law graphs, nearest-neighbor graphs, scale-free graphs, and many others. These model networks with random connections (Erdős-Rényi graphs), networks of brain neurons (small-world graphs), social networks (scale-free graphs), and others.

As in classical signal processing, graph signals can have properties that are specific to graphs and that have not yet been fully understood. They can also be represented via basis atoms and can have a spectral representation. In particular, the graph Fourier transform allows us to develop the intuition gathered in the classical setting and extend it to graphs; we can talk about the notions of frequency and bandlimitedness,

Geert Leus¹, Antonio G. Marques², José M.F. Moura¹, Antonio Ortega³, and David I. Shuman⁴



Digital Object Identifier 10.1109/MSP.2023.3262966
Date of current version: 1 June 2023

75TH ANNIVERSARY OF SIGNAL PROCESSING SOCIETY SPECIAL ISSUE

Graph Signal Processing


History, development, impact, and outlook

Xiaowen Dong, Dorina Thanou, Laura Toni,
Michael Bronstein, and Pascal Frossard

GRAPH SIGNAL PROCESSING:
FOUNDATIONS AND EMERGING DIRECTIONS

Graph Signal Processing for Machine Learning

A review and new perspectives



The effective representation, processing, analysis, and visualization of large-scale structured data, especially those related to complex domains, such as networks and graphs, are one of the key questions in modern machine learning. Graph signal processing (GSP) is a branch of signal processing models and algorithms that aims at handling data supported on graphs, opens new paths of research to address this challenge. In this article, we review a few important contributions made by GSP concepts and tools, such as graph filters and transforms, to the development of novel machine learning algorithms. In particular, our discussion focuses on the following three aspects: exploiting data structure and relational priors, improving data and computational efficiency, and enhancing model interpretability. Furthermore, we provide new perspectives on the future development of GSP techniques that may serve as a bridge between applied mathematics and signal processing on one side and machine learning and network science on the other. Cross-fertilization across these different disciplines may help unlock the numerous challenges of complex data analysis in the modern age.

Introduction
We live in a connected society. Data collected from large-scale interactive systems, such as biological, social, and financial networks, become largely available. In parallel, the past few decades have seen a significant amount of interest in the machine learning community for network data processing and analysis. Networks have an intrinsic structure that conveys very specific properties to data, e.g., interdependences between data entities in the form of pairwise relationships. These properties are traditionally captured by mathematical representations such as graphs.

In this context, new trends and challenges have been developing fast. Let us consider, for example, a network of interconnected interacting genes and the evolution of individual genes at every point in time. Some typical tasks in network biology related to this type of data are 1) discovery of key genes (via protein grouping) affected by the infection and 2) prediction of how the host organism reacts (in terms of gene expression)

Manuscript received November 25, 2017; revised March 01, 2018; accepted March 23, 2018. First posted April 24, 2018. Corresponding author: Antonio Ortega.

¹ A. Ortega is with the University of Southern California, Los Angeles, CA 90089, USA.
² E-mail: antonio.ortega@usp.ac.edu.

³ A. G. Marques is with EPFL, Lausanne, Switzerland 1015, Lausanne, Switzerland.

⁴ J. M. F. Moura is with Carnegie Mellon University, Pittsburgh, PA 15213, USA.

Digital Object Identifier 10.1109/IEEESP.2018.280126

0885-8129 © 2018 IEEE. Personal use is permitted, but republishing/redistributing requires IEEE permission.

Authorized licensed use limited to: UNIVERSITY AT ALBANY. Downloaded on August 22, 2023 at 19:52:10 UTC from IEEE Xplore. Restrictions apply.
809 PROCEEDINGS OF THE IEEE | Vol. 106, No. 5, May 2018

1053-588X/18/050809-13 \$19.95 © 2018 IEEE. All rights reserved.

Authorized licensed use limited to: UNIVERSITY AT ALBANY. Downloaded on August 22, 2023 at 19:52:10 UTC from IEEE Xplore. Restrictions apply.

49

Digital Object Identifier 10.1109/MSP.2020.3014591
Date of current version: 28 October 2020

Authorized licensed use limited to: UNIVERSITY AT ALBANY. Downloaded on August 22, 2023 at 19:52:10 UTC from IEEE Xplore. Restrictions apply.

117

Learning with covariance matrices: A GSP approach

- Signal and information processing is about exploiting **signal structure**
- **Graph signal processing (GSP):** broaden classical signal processing to graphs



Graph Signal Processing: Overview, Challenges, and Applications

This article presents methods to process data associated to graphs (graph signals) extending techniques (transforms, sampling, and others) that are used for conventional signals.

By ANTONIO ORTEGA¹, Fellow IEEE, PASCAL FROSSARD, Fellow IEEE, JELENA KOVACIĆ, Fellow IEEE, JOSÉ M. F. MOURA², Fellow IEEE, and PIERRE VANDERGHEYNST

ABSTRACT

In this article, we present methods to process data associated to graphs (graph signals) extending

techniques

(transforms, sampling, and others) that are used for conventional signals.

to develop

in domains

and

along with

recently

areas. We

tools, including

Next, we

including

data, and


models,

and graph

Learning with covariance matrices: A GSP approach

- Graph neural networks (GNNs) have been shown to be [Ruiz et al., 2023]
 - stable to (**abstract**) perturbations in graph structure
 - generalizable to graph structures of different sizes
(similar to convolutional neural nets for images)


$$\mathbf{x} = [x_1, \dots, x_6]^T$$



- Covariance matrix is a **data-driven** graph
 - interplay between perturbation theory of covariances and ML over them


Applications to network neuroscience

- Covariance matrices appear commonly in network neuroscience



$$\mathbf{C} = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1m} \\ c_{21} & c_{22} & \cdots & c_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mm} \end{bmatrix}$$

Anatomical covariance matrix




Functional connectome

- Principled ML approaches for **reproducible, transparent, generalizable** findings


Applications to network neuroscience

- Covariance matrices appear commonly in network neuroscience



Anatomical covariance matrix


$$\mathbf{C} = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1m} \\ c_{21} & c_{22} & \cdots & c_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mm} \end{bmatrix}$$



Functional connectome

- Principled ML approaches for **reproducible, transparent, generalizable** findings
- **Brain age gap** is a biomarker that reflects neurodegeneration
 - How VNN theoretical advances provide principled **brain age gap prediction?**

■ Healthy
■ Neurodegeneration



coVariance neural networks

➤ coVariance neural networks (VNNs):

GNNs operating on covariance matrices



Saurabh Sihag^{ID}, Gonzalo Mateos^{ID}, and Alejandro Ribeiro^{ID}

➤ Two tutorial articles in IEEE SPM

- Tutorial article on '*Disentangling neurodegeneration with brain age gap prediction models*' (to appear in 2025)
- Tutorial article on '*CoVariance Neural Networks: Principal Component Analysis Meets Learning with Graphs*' (under preparation)

Disentangling Neurodegeneration With Brain Age Gap Prediction Models

A graph signal processing perspective



Neurodegeneration is the progressive loss of structure or function of neurons in the brain. Reduction in cortical thickness or volume over time has been a workhorse metric used to assess neurodegeneration in clinical settings; see case study 1 in “Case Study 1: Cortical Atrophy Characterizes Neurodegeneration in Alzheimer’s Disease” for a demonstration of cortical atrophy assessment in the context of Alzheimer’s disease (AD) relative to healthy individuals [healthy cohort (HC) group]. Naturally, visual inspection of T1-weighted brain magnetic resonance imaging (MRI) images and associated MRI quantification products are used along with other biological measurements to make a “subjective” assessment about the brain health of an individual. These assessments tend to be subjective because they lack a deterministic relationship between an individual’s health status and the absolute values of the metrics observed within MRI scans [1]. Moreover, such methods cannot adequately account for the statistical complexities inherent within neuroimaging datasets that capture neurodegeneration. In particular, neurodegeneration is a characteristic of the healthy aging process and various neurological disorders [2], exhibiting correlated patterns across brain regions. Such statistical factors motivate well the use of data-driven methods to characterize neurodegeneration.

Automating or improving the analyses of brain MRI images is appealing for several reasons: MRI is a noninvasive proce-

Outline

- PCA and the graph Fourier transform
- Covariance neural networks (VNNs)
- Theory of VNNs: Stability and transferability
- Application: Principled brain age gap prediction with VNNs
- Variants of VNNs


Key takeaways

- VNNs offer a novel GSP-inspired perspective to PCA addressing challenges in modern data analysis
- Principled deep learning solution for finite-data regimes
 - Stability and transferability
- VNNs address methodological/conceptual obscurities in brain age gap prediction

PCA and Graph Fourier Transform


A graph filter implementation of PCA inference

- **To show:** PCA-based inference can be implemented with a polynomial over $\hat{\mathbf{C}}$



A graph filter implementation of PCA inference

- **To show:** PCA-based inference can be implemented with a polynomial over $\hat{\mathbf{C}}$



- **How:** Follows from the graph Fourier transform analysis of $\sum_{k=0}^K h_k \hat{\mathbf{C}}^k$
- **Implications:**
 - Alternative implementation of PCA-based inference using polynomial over $\hat{\mathbf{C}}$
 - But more importantly, polynomial implementation is **stable, transferable**


Preliminaries: Graph

➤ **Graph:** a triplet (V, E, W)

- A set of **nodes** $V = \{1, \dots, m\}$
- A set of (undirected) **edges** $E \subseteq V \times V$

Edge between node i and j denoted by (i, j)

- An **edge function** $W: E \mapsto \mathbb{R}$ that maps edge (i, j) to weight $w_{ij} \in \mathbb{R}$



Preliminaries: Graph

➤ **Graph:** a triplet (V, E, W)


- A set of **nodes** $V = \{1, \dots, m\}$
- A set of (undirected) **edges** $E \subseteq V \times V$

Edge between node i and j denoted by (i, j)

- An **edge function** $W: E \mapsto \mathbb{R}$ that maps edge (i, j) to weight $w_{ij} \in \mathbb{R}$


➤ **Adjacency matrix** representation of graph

$$[\mathbf{A}]_{ij} = \begin{cases} w_{ij}, & \text{if } (i, j) \in E, \\ 0, & \text{otherwise} \end{cases}$$



Preliminaries: Graph signal

- **Graph signals** are mappings $x: V \mapsto \mathbb{R}$
 - ➡ graph signal is defined on the vertices of the graph
- **Graph signal** can be represented as a vector $\mathbf{x} \in \mathbb{R}^m$
 - ➡ x_i denotes the graph signal at i -th vertex in V




[Shuman, 2013]

$$\mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_9 \end{bmatrix} = \begin{bmatrix} 0.6 \\ 0.7 \\ 0.3 \\ \vdots \\ 0.7 \end{bmatrix}$$

Preliminaries: Graph shift operator (GSO)

- To understand and analyze graph signal \mathbf{x} , GSP accounts for the graph structure
- Graph structure is encoded in a **graph shift operator** $\mathbf{S} \in \mathbb{R}^{m \times m}$
 - ➡ $[\mathbf{S}]_{ij} = 0$ for $i \neq j$ and $(i, j) \notin E$ (\mathbf{S} captures local graph structure)



$$\mathbf{S} = \begin{pmatrix} S_{11} & S_{12} & 0 & 0 & S_{15} & 0 \\ S_{21} & S_{22} & S_{23} & 0 & S_{25} & 0 \\ 0 & S_{23} & S_{33} & S_{34} & 0 & 0 \\ 0 & 0 & S_{43} & S_{44} & S_{45} & S_{46} \\ S_{51} & S_{52} & 0 & S_{54} & S_{55} & 0 \\ 0 & 0 & 0 & S_{64} & 0 & S_{66} \end{pmatrix}$$

- Examples: adjacency matrix, Laplacian

Covariance matrix is a **data-driven** adjacency matrix

Preliminaries: Graph Fourier Transform (GFT)

- Generically, eigendecomposition of GSO $\mathbf{S} = \mathbf{U}\Phi\mathbf{U}^{-1}$
- **GFT** is the projection of graph signal on the eigenvector space \mathbf{U}

$$\tilde{\mathbf{x}} = \mathbf{U}^{-1}\mathbf{x}$$

- **Inverse GFT** is defined as

$$\mathbf{x} = \mathbf{U}\tilde{\mathbf{x}}$$

➡ Eigenvectors $\mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_m]$ are the frequency basis

When GSO is covariance matrix...

- GFT over covariance matrix

Given eigendecomposition

$$\hat{\mathbf{C}} = \hat{\mathbf{V}} \hat{\Lambda} \hat{\mathbf{V}}^T$$

GFT of \mathbf{x} is

$$\tilde{\mathbf{x}} = \hat{\mathbf{V}}^T \mathbf{x}$$

When GSO is covariance matrix...

- GFT over covariance matrix

Given eigendecomposition

$$\hat{\mathbf{C}} = \hat{\mathbf{V}} \hat{\Lambda} \hat{\mathbf{V}}^T$$

GFT of \mathbf{x} is

$$\tilde{\mathbf{x}} = \hat{\mathbf{V}}^T \mathbf{x}$$

PCA transform is GFT with respect to the covariance graph!

- PCA transform

Projection of sample \mathbf{x} on principal components of $\hat{\mathbf{C}}$

PCA transform: $\tilde{\mathbf{x}} = \hat{\mathbf{V}}^T \mathbf{x}$




Preliminaries: Graph filter

- **Graph filter \mathbf{H}** maps graph signal \mathbf{x} to another graph signal \mathbf{z} via linear-shift-and-sum operation

$$\mathbf{z} = \mathbf{H}(\mathbf{S})\mathbf{x},$$

$$\text{where } \mathbf{H} := h_0 \mathbf{S}^0 + h_1 \mathbf{S}^1 + h_2 \mathbf{S}^2 + \cdots + h_K \mathbf{S}^K = \sum_{k=0}^K h_k \mathbf{S}^k$$





[Isufi et. al, IEEE TSP, 2024]

Graph filter on covariance matrix

- Covariance matrix forms a fully-connected graph where
 - nodes are features
 - edges are covariance values
- Graph filter on covariance matrix $\hat{\mathbf{C}}$ is defined as

$$\mathbf{H}(\hat{\mathbf{C}}) = \sum_{k=0}^K h_k \hat{\mathbf{C}}^k \mathbf{x}$$



CoVariance filter

- Analogy between $\mathbf{H}(\hat{\mathbf{C}})$ and PCA
 - Using eigendecomposition $\hat{\mathbf{C}} = \hat{\mathbf{V}}\hat{\Lambda}\hat{\mathbf{V}}^T$, it follows that

$$\mathbf{z} = \mathbf{H}(\hat{\mathbf{C}})\mathbf{x} = \sum_{k=0}^K h_k \hat{\mathbf{C}}^k \mathbf{x} = \sum_{k=0}^K h_k \hat{\mathbf{V}} \hat{\Lambda}^k \hat{\mathbf{V}}^T \mathbf{x} = \hat{\mathbf{V}} \left(\underbrace{\sum_{k=0}^K h_k \hat{\Lambda}^k}_{\text{Frequency response}} \right) \hat{\mathbf{V}}^T \mathbf{x}$$



CoVariance filter

- Analogy between $\mathbf{H}(\hat{\mathbf{C}})$ and PCA

- Using eigendecomposition $\hat{\mathbf{C}} = \hat{\mathbf{V}}\hat{\Lambda}\hat{\mathbf{V}}^T$, it follows that

$$\mathbf{z} = \mathbf{H}(\hat{\mathbf{C}})\mathbf{x} = \sum_{k=0}^K h_k \hat{\mathbf{C}}^k \mathbf{x} = \sum_{k=0}^K h_k \hat{\mathbf{V}} \hat{\Lambda}^k \hat{\mathbf{V}}^T \mathbf{x} = \hat{\mathbf{V}} \left(\underbrace{\sum_{k=0}^K h_k \hat{\Lambda}^k}_{\text{Frequency response}} \right) \hat{\mathbf{V}}^T \mathbf{x}$$

PCA


- GFT of coVariance filter output \mathbf{z} and PCA are **equivalent**

$$\tilde{\mathbf{z}} = \left(\sum_{k=0}^K h_k \hat{\Lambda}^k \right) \hat{\mathbf{V}}^T \mathbf{x}$$

i -th component of $\tilde{\mathbf{z}}$ is modulated by $h(\lambda_i) = \sum_{k=0}^K h_k \lambda_i^k$


Learning with coVariance filter versus PCA-based learning

- Learning with a **coVariance** filter





Learning with coVariance filter versus PCA-based learning

- Learning with a **coVariance** filter





Learning with coVariance filter versus PCA-based learning

- Learning with a **coVariance** filter





Learning with coVariance filter versus PCA-based learning

- Learning with a **coVariance** filter





- PCA-based learning





Learning with coVariance filter versus PCA-based learning

- Learning with a **coVariance** filter




- **PCA**-based learning



coVariance Neural Networks (VNNs)

coVariance filters as convolutional operators

- Operation $\hat{\mathbf{C}}^k \mathbf{x}$ performs a k -shift of signal \mathbf{x} over graph defined by $\hat{\mathbf{C}}$




- Parameters $\{h_k\}$ are called **filter taps**, are **scalars** and **learnable** parameters

CoVariance Neural Networks (VNNs)

- coVariance filters can learn only **linear** representations
- To accommodate learn **non-linear** representations, concatenate coVariance filter with pointwise non-linearity σ (for e.g., ReLU, sigmoid, etc.)


CoVariance Neural Networks (VNNs)

- coVariance filters can learn only **linear** representations
- To accommodate learn **non-linear** representations, concatenate coVariance filter with pointwise non-linearity σ (for e.g., ReLU, sigmoid, etc.)




CoVariance Neural Networks (VNNs)

- coVariance filters can learn only **linear** representations
- To accommodate learn **non-linear** representations, concatenate coVariance filter with pointwise non-linearity σ (for e.g., ReLU, sigmoid, etc.)




VNNs outperform PCA (regression task)

Synthetic data
(Friedman regression problem)



Neuroimaging data
(age prediction task)




Covariance Filters and Neural Networks

Covariance filters


- A covariance filter is a **polynomial in the covariance matrix $\hat{\mathbf{C}}$**

$$\mathbf{H}(\hat{\mathbf{C}}) = \sum_{k=0}^K h_k \hat{\mathbf{C}}^k \mathbf{x}$$

- We train the filter coefficients h_k to accomplish some task




CoVariance Neural Networks (VNNs)



- A VNN is a composition of layers
- Each of which is a composition of
 - ... a **covariance filter**
 - ... with a **pointwise nonlinearity**
- $\Phi(\mathbf{x}; \hat{\mathbf{C}}, \mathcal{H})$ represents VNN output
- \mathcal{H} is the set of trainable filter taps

Covariance Filters are Implicitly Equivalent to PCA




Covariance Filters are Implicitly Equivalent to PCA

- The difference is that covariance filters (and VNNs) **do not require eigenvectors**

Stability: Leading to more stable signal processing

Transferability: And the possibility of transferring trained filters across scales




Stable Inference with VNNs

Stability of inference with PCA and VNNs

- PCA-driven inference can be **unstable**
 - stochastic perturbations due to **finite sample effect**
- VNNs provide **stable** outcomes
 - enhanced reproducibility
 - avoid overfitting

Performance on regression task



\hat{C}_n : estimated from n samples

Stochastic perturbations in sample covariance matrix

- Recall: Sample covariance matrix $\hat{\mathbf{C}}$ is estimate of true covariance matrix \mathbf{C}

$$\hat{\mathbf{C}} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T \quad \mathbf{C} = \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T]$$

➡ eigenvectors/eigenvalues $\hat{\mathbf{V}}, \hat{\boldsymbol{\Lambda}}$ of $\hat{\mathbf{C}}$ are estimates of $\mathbf{V}, \boldsymbol{\Lambda}$ of \mathbf{C}

Stochastic perturbations in sample covariance matrix

- Recall: Sample covariance matrix $\hat{\mathbf{C}}$ is estimate of true covariance matrix \mathbf{C}

$$\hat{\mathbf{C}} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T \quad \mathbf{C} = \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T]$$

➡ eigenvectors/eigenvalues $\hat{\mathbf{V}}, \hat{\boldsymbol{\Lambda}}$ of $\hat{\mathbf{C}}$ are estimates of $\mathbf{V}, \boldsymbol{\Lambda}$ of \mathbf{C}

- Convergence between $\hat{\mathbf{V}}, \hat{\boldsymbol{\Lambda}}$ and $\mathbf{V}, \boldsymbol{\Lambda}$ [*]

$$\|\hat{\mathbf{V}}_{\mathbf{x}} - \mathbf{V}_{\mathbf{x}}\| = \mathcal{O}\left(\frac{1}{n^{1/2} \min_{i \neq j} |\lambda_i - \lambda_j|}\right)$$

[*] Loukas, Andreas, 2017

Stochastic perturbations in sample covariance matrix

- Recall: Sample covariance matrix $\hat{\mathbf{C}}$ is estimate of true covariance matrix \mathbf{C}

$$\hat{\mathbf{C}} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T \quad \mathbf{C} = \mathbb{E}[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T]$$

➡ eigenvectors/eigenvalues $\hat{\mathbf{V}}, \hat{\boldsymbol{\Lambda}}$ of $\hat{\mathbf{C}}$ are estimates of $\mathbf{V}, \boldsymbol{\Lambda}$ of \mathbf{C}

- Convergence between $\hat{\mathbf{V}}, \hat{\boldsymbol{\Lambda}}$ and $\mathbf{V}, \boldsymbol{\Lambda}$ [*]

$$\|\hat{\mathbf{V}}_{\mathbf{x}} - \mathbf{V}_{\mathbf{x}}\| = \mathcal{O}\left(\frac{1}{n^{1/2} \min_{i \neq j} |\lambda_i - \lambda_j|}\right)$$

➡ **Unstable PCA transform when eigenvalues of covariance are close**

[*] Loukas, Andreas, 2017

Stability of coVariance filter

- How to gauge stability?

$$\mathbf{x} \longrightarrow \boxed{\mathbf{H}(\hat{\mathbf{C}})} \longrightarrow \mathbf{z} = \mathbf{H}(\hat{\mathbf{C}})\mathbf{x}$$
$$\hat{\mathbf{C}} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^\top$$

➡ Output \mathbf{z} must be robust to number of samples n used to estimate $\hat{\mathbf{C}}$

Stability of coVariance filter

- How to gauge stability?

$$\mathbf{x} \longrightarrow \boxed{\mathbf{H}(\hat{\mathbf{C}})} \longrightarrow \mathbf{z} = \mathbf{H}(\hat{\mathbf{C}})\mathbf{x} \quad \hat{\mathbf{C}} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^\top$$


➡ Output \mathbf{z} must be robust to number of samples n used to estimate $\hat{\mathbf{C}}$

- Compare filter outputs for sample and true covariance matrix

$$\mathbf{x} \longrightarrow \boxed{\mathbf{H}(\hat{\mathbf{C}})} \longrightarrow \mathbf{z} = \mathbf{H}(\hat{\mathbf{C}})\mathbf{x} \quad \mathbf{x} \longrightarrow \boxed{\mathbf{H}(\mathbf{C})} \longrightarrow \mathbf{z} = \mathbf{H}(\mathbf{C})\mathbf{x}$$

➡ metric of interest: $\|\mathbf{H}(\hat{\mathbf{C}}) - \mathbf{H}(\mathbf{C})\|$

Stability of coVariance filter




Stability result [Sihag et al., 2022]

$$\left\| H(\hat{C}) - H(C) \right\| = \mathcal{O} \left(\frac{1}{n^{1/2-\varepsilon}} \right)$$

} coVariance filter output is asymptotically consistent

Stability of coVariance filter



Stability result [Sihag et al., 2022]

$$\left\| H(\hat{C}) - H(C) \right\| = \mathcal{O} \left(\frac{1}{n^{1/2-\varepsilon}} \right)$$

Assumption.



Frequency response of filter $H(C)$ satisfies

$$|h(\lambda_i) - h(\lambda_j)| \leq Q \frac{|\lambda_i - \lambda_j|}{k_i}$$



- } coVariance filter output is asymptotically consistent
- } coVariance filter sacrifices discriminability between close eigenvalues for stability

Recall: Learning with coVariance filter versus PCA-based learning

➤ Learning with a **coVariance** filter





➤ PCA-based learning





Why is coVariance filter more stable than PCA?

➤ Learning with a **coVariance** filter





➤ PCA-based learning



Why is coVariance filter more stable than PCA?



➤ Learning with a **coVariance** filter



Eigenvalue,
Eigenvector

$\hat{\lambda}_4, \hat{\lambda}_3,$
 $\hat{v}_4 \hat{v}_3$

➤ PCA-based learning



Eigenvalue,
Eigenvector

$\hat{\lambda}_4, \hat{\lambda}_3,$
 $\hat{v}_4 \hat{v}_3$

Overfitting on the ordering of eigenvalues is source of instability

Stability of VNNs


- VNNs inherit the stability from coVariance filters
 - Stability bound depends on the bound for filters

$$\left\| \mathbf{H}(\hat{\mathbf{C}}) - \mathbf{H}(\mathbf{C}) \right\| = \mathcal{O} \left(\frac{1}{n^{\frac{1}{2}} - \varepsilon} \right) = \alpha_n$$

- For a VNN with L layers and F filters in parallel,


$$\left\| \Phi(\mathbf{x}, \hat{\mathbf{C}}; \mathcal{H}) - \Phi(\mathbf{x}, \mathbf{C}; \mathcal{H}) \right\| \leq LF^{L-1} \alpha_n$$

- Stability bound increases with number of layers and size of filter banks




Stability of VNNs: Experiments

- Regression task



Cortical thickness
data




→ Estimate of age

- Comparison against PCA-regression

Data: cortical thickness dataset ($m = 104$) from ($n = 341$) human subjects

- **Metric:** MAE (mean absolute error)



VNN: coVariance Neural Network

PCA-LR: PCA-regression with linear kernel


PCA-rbf: PCA regression with rbf kernel

VNN outperforms PCA and is more stable

Transferability of VNNs

Empirical evidence of transferability across multiscale data

- Transferability across multiscale datasets
 - **Multiscale** datasets capture same phenomenon at different scales




Transferability across datasets with different number of features


Training	Testing	
	100-feature dataset	300-feature dataset
100-feature dataset	5.39 ± 0.084	5.5 ± 0.101

Transferability

- Learning models could generalize to **compatible** datasets
- **compatible:** different dimensionalities and describing the same domain




Brain imaging data




Credit: Mustafa Aksoy, UAlbany

Transferability

- Learning models could generalize to **compatible** datasets
- **compatible:** different dimensionalities and describing the same domain



Brain imaging data





Credit: Mustafa Aksoy, UAlbany

- **Motivation:** novel metric for generalizability, managing high dimensional data...

Transferability


- Most statistical approaches, including PCA, operate within the dimensionality
 - ➡ seamless transference not possible across different dimensionalities
- **This section: How do VNNs transfer?**

When is transference successful?



Credit: Mustafa Aksoy, UAlbany

coVariance filters are scale-free models




- A coVariance filter $\mathbf{H}(\cdot)$ with scalar filter taps $\{h_k\}$ can process dataset (covariance matrix) of any arbitrary dimensionality: **scale-free model**


VNNs as scale-free models



m_1 -dimensional data processing



m_2 -dimensional data processing




learnable parameters

How to compare $\Phi(\mathbf{x}_{m_1}; \mathbf{C}_{m_1}, \mathcal{H})$ and $\Phi(\mathbf{x}_{m_2}; \mathbf{C}_{m_2}, \mathcal{H})$?


VNNs as scale-free models



m_1 -dimensional data processing



data processing in continuous limit



learnable parameters


Continuous limit of covariance matrices

as $m \rightarrow \infty$


How to compare $\Phi(\mathbf{x}_{m_1}; \mathbf{C}_{m_1}, \mathcal{H})$ and $\Phi(y_{\mathbf{x}}; \mathbf{W}, \mathcal{H})$?

Graphons as continuous limits

- Graphs can have **limit objects** with uncountable number of nodes
- Example: Stochastic block models [Ruiz et al., 2021]



Balanced SBM



Unbalanced SBM

Graphons as continuous limits

➤ **Graphon:** A graphon is a symmetric, bounded measurable function

- Node labels are graphon arguments $u \in [0,1]$
- edge weights are graphon values $\mathbf{W}(u, v) = \mathbf{W}(v, u)$

$$\mathbf{W} : [0, 1]^2 \mapsto \mathbb{R}$$


Graphons as continuous limits

➤ **Graphon:** A graphon is a symmetric, bounded measurable function


- Node labels are graphon arguments $u \in [0,1]$
- edge weights are graphon values $\mathbf{W}(u, v) = \mathbf{W}(v, u)$

$$\mathbf{W} : [0, 1]^2 \mapsto \mathbb{R}$$


➤ Transferability when covariance matrix is part of some converging sequence




Redefining VNNs in continuous domain




Redefining VNNs in continuous domain




Redefining VNNs in continuous domain



Problem formulation for transferability




Problem formulation for transferability



Find ϑ , such that, $\|y_{m_1} - y\|_2 \leq \vartheta$


VNNs are provably transferable



Transferability bound* [Sihag et al., 2024]

$$\|y_{m_1} - y\| \propto \mathcal{O}\left(\frac{1}{m_1^{3\zeta/2-1}}\right), \text{ for } \zeta \in (2/3, 1]$$


VNNs are provably transferable




Transferability bound* [Sihag et al., 2024]

$$\|y_{m_1} - y\| \propto \mathcal{O}\left(\frac{1}{m_1^{3\zeta/2-1}}\right), \text{ for } \zeta \in (2/3, 1]$$

*Assumption: data is a discretization of a common continuous model



VNNs are provably transferable




Transferability bound

$$\|y_{m_1} - y_{m_2}\| \propto \mathcal{O}\left(\frac{1}{m_1^{3\zeta/2-1}} + \frac{1}{m_2^{3\zeta/2-1}}\right), \text{ for } \zeta \in (2/3, 1]$$

Experiments

Objective: Brain age gap prediction in HC (healthy) and AD+ (Alzheimer's) cohorts from VNNs trained on 100-feature dataset [Sihag et al., NeurIPS, 2024, JSTSP 2024, SPM 2025]



- ROIs contributing to elevated brain age gap in AD+ across different resolutions
- Brain age gap is elevated in AD+ w.r.t HC cohort in 100-feature dataset
- Results on brain age gap retained after transferring VNN to 300 and 500-feature datasets


Principled brain age gap prediction with VNNs

Brain age gap

- Individual rate of “aging” is different from chronological rate of aging
 - Driven by environment, genetics, **neurodegeneration**
- **Brain age** provides a biological estimate brain age, derived from **neuroimaging**

Brain age gap

- Individual rate of “aging” is different from chronological rate of aging
 - Driven by environment, genetics, **neurodegeneration**
- **Brain age** provides a biological estimate brain age, derived from **neuroimaging**
- The **brain age gap** is the **deviation** between brain age and chronological age




individual risks for neurological,
Brain age gap \propto neuropsychiatric
and neurodegenerative diseases

Neurodegeneration (in terms of cortical atrophy)

- Neurodegeneration is **accelerated decline** of structure or function of the brain
- **Cortical atrophy:** reduction in cortical **thickness/volume/area**

(characteristic of healthy aging and disorders like Alzheimer's disease (AD), frontotemporal dementia (FTD), etc.)




MRI scan

Preprocessing
(e.g., Freesurfer)



$$\mathbf{x} = [x_1, \dots, x_m]$$




Cortical thickness features

x_i is cortical thickness for brain region i

Neuroimaging Data: Basics

- Data sample corresponds to measurement associated with brain (cortical) surface

$$\mathbf{x} = [x_1, \dots, x_m]$$




Anatomic features

- Brain surface is divided according to **brain atlases**
 - ➡ datasets may have **distinct** dimensionalities

- **Multi-resolution** brain atlas discretizes brain surface at multiple resolutions
(for e.g., Schaefer's atlas has resolutions 100-1000)


Case study (Neurodegeneration)

- **Data:** cortical thickness from 3 cohorts
 - HC (healthy)
 - MCI (Mild cognitive impairment)
 - AD (Alzheimer's disease)
- Larger **cortical atrophy** is feature of AD
- MCI is precursor to AD
 - ➡ shows intermediate cortical atrophy between HC and AD
- **Aging** also leads to cortical atrophy




Brain age gap evaluation using ML

Step 1. Train ML model to predict chronological age for healthy controls from cortical thickness features




Brain age gap evaluation using ML

Step 1. Train ML model to predict chronological age for healthy controls from cortical thickness features




Step 2. Linear regression-based age-bias correct for outputs of ML model

Step 3. Obtain **brain age gap** for healthy controls and individuals with neurodegenerative condition.




Brain age gap prediction is a transfer learning problem

- Train ML model to predict age on a **large dataset** (healthy population)




- Apply the **pre-trained** ML model on a **target dataset** (neurodegeneration)



- Brain age gap is the **residual** of the model


Brain age gap prediction is a transfer learning problem

- **Some observations about a meaningful brain age gap**
 - We expect model performance to **degrade** in **target population**
 - ✓ Degradation in performance (residuals) must be in a **specific direction**
 - ✓ Degradation in performance (residuals) \propto **disease severity/status**




Choice of learning parametrization

- Choice of ML model dictates how data is leveraged to gauge brain age gap
- Prevalent approaches focus on achieving perfect pre-training performance
 - **Performance-driven approaches**
- **Performance-driven approaches do not guarantee ‘meaningful’ brain age gap**



Choice of learning parametrization

- Neural networks are prevalent in performance-driven approaches
- A Neural Network may **not be interpretable** and prone to **overfitting**



Choice of learning parametrization

- Neural networks are prevalent in performance-driven approaches
- A Neural Network may **not be interpretable** and prone to **overfitting**
 - ➡ **methodological obscurity** in brain age gap prediction pipeline


Performance in **pre-training** does not dictate **meaningful residuals** in **target population**

A principled approach to brain age gap prediction


- **Focus on residuals** of the ML model, not prediction performance
- **Qualitative evaluation** during pre-training
 - what does the model learn during **pre-training** on **healthy population**?
- **Interpretability/explainability:**
 - what's driving elevated brain age gap (residuals) in **neurodegeneration**?
- **Generalizability** to diverse target populations

Sihag et al., 2025 (SPM, to appear)


VNNs provide an anatomically interpretable and explainable brain age gap



VNNs provide an anatomically interpretable and explainable brain age gap




VNNs provide an anatomically interpretable and explainable brain age gap




Network neuroscience

➤ Modeling brain as a network (**connectomes**)



$$C = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1m} \\ c_{21} & c_{22} & \cdots & c_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mm} \end{bmatrix}$$

Anatomical covariance matrix
(structural connectome)




Functional connectome

➤ Motivation

- Significant redundancies in brain structural/functional features
- Brain structure/function is compromised in neurodegeneration


Covariance matrices in network neuroscience

- Covariance matrices appear commonly in network neuroscience



$$\mathbf{C} = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1m} \\ c_{21} & c_{22} & \cdots & c_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mm} \end{bmatrix}$$


Anatomical covariance matrix



Functional connectome


Covariance matrices in network neuroscience

- Covariance matrices appear commonly in network neuroscience



Anatomical covariance matrix


$$\mathbf{C} = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1m} \\ c_{21} & c_{22} & \cdots & c_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mm} \end{bmatrix}$$



Functional connectome

- Inference over covariance matrices in network neuroscience

- **Traditional** statistical approaches (for e.g., PCA)
 - Interpretable, suitable for low data regimes
- **Deep learning** approaches (for e.g., GNNs)
 - Enhanced expressivity, improved performance




VNNs are well suited for neuroimaging data analysis

- Properties of VNNs make them appealing for neuroimaging data analysis
- **Connections with PCA** → transparent outcomes by leveraging spectrum of covariance matrix
- **Stability** → reproducible outcomes in limited data settings
- **Transferability** → enhanced generalizability and robustness to choice of brain atlases

Anatomical covariance matrix as a graph



- Covariance matrix is a **data-driven** graph

$$\mathbf{x} = [x_1, \dots, x_6]^T$$



Covariance matrix as a fully-connected graph

$$\hat{\mathbf{C}} = \frac{1}{n-1} \sum_{i=1}^n (\mathbf{x}_i - \hat{\boldsymbol{\mu}})(\mathbf{x}_i - \hat{\boldsymbol{\mu}})^T, \text{ where } \hat{\boldsymbol{\mu}} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$$




$$\mathbf{C} = \begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1m} \\ c_{21} & c_{22} & \cdots & c_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mm} \end{bmatrix}$$


Anatomical covariance matrix
(estimated from cortical features)

VNN vs PCA on age prediction task

- Regression task



Cortical thickness
data




→ Estimate of age

- Comparison against PCA-regression

Data: cortical thickness dataset ($m = 104$) from ($n = 341$) human subjects

- **Metric:** MAE (mean absolute error)




VNN: coVariance Neural Network

PCA-LR: PCA-regression with linear kernel


PCA-rbf: PCA regression with rbf kernel

VNN outperforms PCA and is **more stable**
[Sihag et al., 2022]


VNNs provide an anatomically interpretable and explainable brain age gap




VNNs provide an anatomically interpretable and explainable brain age gap



VNNs provide an anatomically interpretable and explainable brain age gap



VNNs provide an anatomically interpretable and explainable brain age gap




Experiments

- Participants from OASIS-3 dataset, 148 cortical thickness features per individual
(Distrieux brain atlas)


	HC	AD
Number	611	194
Age	68.38 (7.62)	74.72 (7.02)
Sex (m/f)	260/351	100/94
CDR sum of boxes	0	3.45 (1.74)

HC group: cognitively normal
AD group: AD diagnosis
CDR: Clinical dementia rating

- Brain age gap is elevated in **AD** group and correlated with CDR sum of boxes





Anatomical interpretability



Experiments

- VNN **distinctly** exploits eigenvectors in AD and HC groups




➡ explains anatomical interpretability of brain age gap in AD


Experiments

- Whole brain cortical thickness dataset for Frontotemporal Dementia (FTD)
 - Healthy controls (HC, n = 114, age = 64.51 ± 6.51 years, 65 females)
 - FTD diagnosis (FTD, n = 119, age = 64.72 ± 6.78 years, 47 females)
- 68 cortical thickness features (Desikan-Killiany atlas)


Brain age gap distributions




Anatomic interpretability



Explaining anatomic interpretability



VNNs are provably transferable









Transferability bound

$$\|y_{m_1} - y_{m_2}\| \propto \mathcal{O}\left(\frac{1}{m_1^{3\zeta/2-1}} + \frac{1}{m_2^{3\zeta/2-1}}\right), \text{ for } \zeta \in (2/3, 1]$$

Recap: Transferability of VNNs cross-validates brain age gap in multi-resolution setting


Objective: Brain age gap prediction in **HC (healthy)** and **AD+ (Alzheimer's)** cohorts from VNNs trained on 100-feature dataset




- ROIs contributing to elevated brain age gap in **AD+** across different resolutions
- Brain age gap is elevated in **AD+** w.r.t **HC** cohort in 100-feature dataset
- Results on brain age gap retained after transferring VNN to 300 and 500-feature datasets

VNNs as pre-trained models...

- Uncovering **disease heterogeneity** with VNNs as pre-trained models



- VNNs offer more **significant clinical stratification** than raw anatomical features




VNN enhances the
clinical relevance of
anatomical features

Alignment of covariance with learning: An NTK perspective


Are covariance matrices suitable for a learning task?

- Covariance matrices add a **meaningful** inductive bias to neural nets
 - Covariance matrix captures the (linear) structure
- VNNs provide the bridge between PCA and GNNs




Are covariance matrices suitable for a learning task?

- Covariance matrices add a **meaningful** inductive bias to neural nets
 - Covariance matrix captures the (linear) structure
- VNNs provide the bridge between PCA and GNNs
- **Can we quantify the suitability of covariance to learning objective?**
 - Is performance good?
 - good generalization?



Are covariance matrices suitable for a learning task?

- Covariance matrices add a **meaningful** inductive bias to neural nets
 - Covariance matrix captures the (linear) structure
- VNNs provide the bridge between PCA and GNNs
- **Can we quantify the suitability of covariance to learning objective?**
 - Is performance good?
 - good generalization?
- **Neural tangent kernels (NTKs)-driven insights for VNNs**




Neural tangent kernels (NTK)

- NTKs describe the **evolution of neural nets** during training by **gradient descent**

- **Example:**


Predict \mathbf{y} from \mathbf{x} using
ML model \mathbf{f} (parameters \mathbf{h})



Neural tangent kernels (NTK)

- NTKs describe the **evolution of neural nets** during training by **gradient descent**
- **Example:**

Predict \mathbf{y} from \mathbf{x} using
ML model \mathbf{f} (parameters \mathbf{h})




Loss : $\mathcal{L} = \text{mean squared error (MSE)}$

$$= \sum_{i \in \text{Data}} \|\mathbf{y}_i - \mathbf{f}(\mathbf{x}_i, \mathbf{h})\|^2$$

- Optimize parameters \mathbf{h} using **gradient descent**

Neural tangent kernels (NTK)

- **Evolution of gradient descent**
(linear approximation)



$$\mathbf{f}(\mathbf{x}, \mathbf{h}^{(t+1)}) = \mathbf{f}(\mathbf{x}, \mathbf{h}^{(t)}) - \eta \Theta(\mathbf{x}, \mathbf{h}^{(t)}) \cdot (\mathbf{f}(\mathbf{x}, \mathbf{h}^{(t)}) - y)$$

η : learning rate

$\Theta(\mathbf{x}, \mathbf{h}^{(t)})$: NTK matrix

$$\Theta(\mathbf{x}_i, \mathbf{x}_j) = \nabla_{\mathbf{h}} \mathbf{f}(\mathbf{x}, \mathbf{h})^{\top} \nabla_{\mathbf{h}} \mathbf{f}(\mathbf{x}, \mathbf{h})$$

Depends on model
architecture + input

Neural tangent kernels (NTK)

- **Evolution of gradient descent**

(linear approximation)

$$\mathbf{f}(\mathbf{x}, \mathbf{h}^{(t+1)}) = \mathbf{f}(\mathbf{x}, \mathbf{h}^{(t)}) - \eta \Theta(\mathbf{x}, \mathbf{h}^{(t)}) \cdot (\mathbf{f}(\mathbf{x}, \mathbf{h}^{(t)}) - y)$$

η : learning rate


$\Theta(\mathbf{x}, \mathbf{h}^{(t)})$: NTK matrix

Depends on model architecture + input

$$\Theta(\mathbf{x}_i, \mathbf{x}_j) = \nabla_{\mathbf{h}} \mathbf{f}(\mathbf{x}, \mathbf{h})^{\top} \nabla_{\mathbf{h}} \mathbf{f}(\mathbf{x}, \mathbf{h})$$

- For neural networks with **infinite width**, NTK matrix is constant w.r.t \mathbf{h}

$$\Theta(\mathbf{x}, \mathbf{h}^{(t)}) \rightarrow \Theta(\mathbf{x})$$



Neural tangent kernels (NTK)

- **Convergence of gradient descent dictated by alignment between NTK and data**

$$\|f(x, h^{(t)}) - y\|_2^2 \propto \mathcal{A} \quad \text{where} \quad \mathcal{A} = y^\top \Theta y$$

Larger \mathcal{A}  faster convergence


Neural tangent kernels (NTK)

- **Convergence of gradient descent dictated by alignment between NTK and data**

$$\|\mathbf{f}(\mathbf{x}, \mathbf{h}^{(t)}) - \mathbf{y}\|_2^2 \propto -\mathcal{A} \quad \text{where} \quad \mathcal{A} = \mathbf{y}^\top \Theta \mathbf{y}$$

Larger \mathcal{A} \rightarrow faster convergence

Example




$$\mathcal{A}_1 < \mathcal{A}_2 < \mathcal{A}_3 < \mathcal{A}_4$$

Larger alignment implies better convergence/ loss

NTK for covariance filter

- For a covariance filter $\mathbf{H}(\hat{\mathbf{C}}) = \sum_{k=0}^K h_k \hat{\mathbf{C}}^k$,


- NTK is $\Theta = \sum_{k=0}^K \hat{\mathbf{C}}^k \mathbf{x} \mathbf{x}^\top \hat{\mathbf{C}}^k$



Alignment for covariance filter

- For a covariance filter $\mathbf{H}(\hat{\mathbf{C}}) = \sum_{k=0}^K h_k \hat{\mathbf{C}}^k$,

- NTK is $\Theta = \sum_{k=0}^K \hat{\mathbf{C}}^k \mathbf{x} \mathbf{x}^\top \hat{\mathbf{C}}^k$



→ convergence of learning with covariance filter is dictated by

$$\mathcal{A} = \mathbf{y}^\top \left(\sum_{k=0}^K \hat{\mathbf{C}}^k \mathbf{x} \mathbf{x}^\top \hat{\mathbf{C}}^k \right) \mathbf{y}$$

Alignment between \mathbf{x} , $\hat{\mathbf{C}}$, and \mathbf{y}

Khalafi et al., 2024


Data-driven graph by optimizing alignment

- Treating alignment as an optimization objective
 - **Goal:** find the *optimal* graph shift operator

$$\begin{aligned}\mathbf{S}^* &= \max_{\mathbf{S}} \mathcal{A}(\mathbf{S}) \\ &= \max_{\mathbf{S}} \mathbf{y}^\top \left(\sum_{k=0}^K \mathbf{S}^k \mathbf{x} \mathbf{x}^\top \mathbf{S}^k \right) \mathbf{y}\end{aligned}$$

➡ Find \mathbf{S} that maximizes

$$\sum_{k=0}^K (\mathbf{y}^\top \mathbf{S} \mathbf{x})^2$$



Khalafi et al., 2024


Data-driven graph by optimizing alignment

- Treating alignment as an optimization objective
 - **Goal:** find the *optimal* graph shift operator

$$\begin{aligned}\mathbf{S}^* &= \max_{\mathbf{S}} \mathcal{A}(\mathbf{S}) \\ &= \max_{\mathbf{S}} \mathbf{y}^T \left(\sum_{k=0}^K \mathbf{S}^k \mathbf{x} \mathbf{x}^T \mathbf{S}^k \right) \mathbf{y}\end{aligned}$$

➡ Find \mathbf{S} that maximizes

$$\sum_{k=0}^K (\mathbf{y}^T \mathbf{S}^k \mathbf{x})^2$$



- Correlation between
- Graph shift operator \mathbf{S}
 - Input \mathbf{x}
 - Output \mathbf{y}

Khalafi et al., 2024

Data-driven covariance graph by optimizing alignment

- **Cross-covariance** graph optimizes alignment

$$\mathbf{S}^* = \frac{1}{2}(\mathbf{x}\mathbf{y}^\top + \mathbf{y}\mathbf{x}^\top)$$

Khalafi et al., 2024


Data-driven covariance graph by optimizing alignment

- **Cross-covariance** graph optimizes alignment

$$\mathbf{S}^* = \frac{1}{2}(\mathbf{x}\mathbf{y}^\top + \mathbf{y}\mathbf{x}^\top)$$

- **Numerical results**

- Time series forecasting: predicting next time step



- Cross-covariance achieves better loss
- GNN with cross-covariance outperforms VNN

Khalafi et al., 2024

Variants of VNNs

Are VNNs enough?

- **Limitations of VNNs**
 - Sample covariance could be poor quality in **low data, high dimensionality setting**
 - High **computational cost** (quadratic in size for dense covariance)
 - No considerations of **temporal, evolving** data
 - Prone to **undesired bias** within the data

Low data, high dimensional settings

- Sample covariance matrix is dense
 - ➡ **noisy** entries in low data, high dimensional settings
 - ➡ computationally inefficient VNNs (quadratic complexity)

Low data, high dimensional settings

- Sample covariance matrix is dense
 - ➡ **noisy** entries in low data, high dimensional settings
 - ➡ computationally inefficient VNNs (quadratic complexity)
- **Solution:** sparsify the sample covariance matrix
 - If **true covariance is sparse**:
 - Improve estimation quality
 - Common in real world
(brain imaging, finance, etc.)


Low data, high dimensional settings

- Sample covariance matrix is dense
 - ➡ **noisy** entries in low data, high dimensional settings
 - ➡ computationally inefficient VNNs (quadratic complexity)
- **Solution:** sparsify the sample covariance matrix
 - If **true covariance is sparse**:
 - Improve estimation quality
 - Common in real world
(brain imaging, finance, etc.)
 - For **generic covariance**:
 - Improve computational efficiency



Sparse VNNs

- **Sparse VNNs:** sparsify the covariance matrix with **thresholding** techniques

Sample covariance



Sparsified covariance




[Cavallo et al., 2024]



Sparse VNNs

- **Sparse VNNs:** sparsify the covariance matrix with **thresholding** techniques

Sample covariance



Sparsified covariance



- What thresholding techniques?
- Are sparse VNNs stable?

} questions to address


[Cavallo et al., 2024]

Hard thresholding


➤ Definition

$$\eta(\hat{\mathbf{C}})_{ij} = \hat{c}_{ij} \text{ if } |\hat{c}_{ij}| \geq \tau/\sqrt{n}, 0 \text{ otherwise}$$

Empirical covariance



Hard-thr covariance



[Cavallo et al., 2024]

Hard thresholding

➤ Definition


$$\eta(\hat{\mathbf{C}})_{ij} = \hat{c}_{ij} \text{ if } |\hat{c}_{ij}| \geq \tau/\sqrt{n}, 0 \text{ otherwise}$$

➤ Stability bound


$$\|\mathbf{H}(\hat{\mathbf{C}}_{\text{thr}}) - \mathbf{H}(\hat{\mathbf{C}})\| = \mathcal{O}\left(\frac{c_0}{n^{1/2}}\right)$$

c_0 : number of non-zero elements in $\hat{\mathbf{C}}_{\text{thr}}$

Empirical covariance



Hard-thr covariance



[Cavallo et al., 2024]

Hard thresholding

- **Definition**


$$\eta(\hat{\mathbf{C}})_{ij} = \hat{c}_{ij} \text{ if } |\hat{c}_{ij}| \geq \tau/\sqrt{n}, 0 \text{ otherwise}$$

- **Stability bound**


$$\|\mathbf{H}(\hat{\mathbf{C}}_{\text{thr}}) - \mathbf{H}(\hat{\mathbf{C}})\| = \mathcal{O}\left(\frac{c_0}{n^{1/2}}\right)$$

c_0 : number of non-zero elements in $\hat{\mathbf{C}}_{\text{thr}}$

Empirical covariance



Hard-thr covariance



- Stability bound for S-VNNs is **tighter** than *dense* VNNs


[Cavallo et al., 2024]

Soft thresholding


➤ Definition

$$\eta(\hat{\mathbf{C}})_{ij} = \hat{c}_{ij} - \text{sign}(\hat{c}_{ij})\tau/n \text{ if } |\hat{c}_{ij}| \geq \tau/\sqrt{n}, 0 \text{ otherwise}$$

Empirical covariance



Soft-thr covariance



[Cavallo et al., 2024]

Soft thresholding

➤ Definition


$$\eta(\hat{\mathbf{C}})_{ij} = \hat{c}_{ij} - \text{sign}(\hat{c}_{ij})\tau/n \text{ if } |\hat{c}_{ij}| \geq \tau/\sqrt{n}, 0 \text{ otherwise}$$

➤ Stability bound


$$\|\mathbf{H}(\hat{\mathbf{C}}_{\text{thr}}) - \mathbf{H}(\hat{\mathbf{C}})\| = \mathcal{O}\left(\frac{c_0}{n^{1/2}}\right)$$

c_0 : number of non-zero elements in $\hat{\mathbf{C}}_{\text{thr}}$

Empirical covariance



Soft-thr covariance




➤ Stability bound for S-VNNs is **tighter** than *dense* VNNs

[Cavallo et al., 2024]

Sparse VNNs: Numerical results

- Train VNNs/PCA on one covariance and test on another covariance estimated from less samples (synthetic dataset)




Results

- S-VNN (both soft and hard thresholding) **outperform** PCA and *dense* VNNs
- VNNs **more stable** than PCA

[Cavallo et al., 2024]

Sparse VNNs: Numerical results


- Classification task on real data
- Datasets (from left to right)
 - Brain recordings: Epilepsy and CNI – classify patient condition
 - Human action recognition: MHEALTH and Realdisp – classify action



[Cavallo et al., 2024]

Sparse VNNs: Numerical results


- Classification task on real data
- Datasets (from left to right)
 - Brain recordings: Epilepsy and CNI – classify patient condition
 - Human action recognition: MHEALTH and Realdisp – classify action



- S-VNNs are **faster** and achieve **better performance** than *dense* VNNs

Limitations of VNNs - 2


- Datasets may contain harmful **biases**
 - For e.g., under-represented groups
 - Biased (unfair) performance
 - Fair PCA might be **unstable**



[Cavallo et al., 2025]

Limitations of VNNs - 2


- Datasets may contain harmful **biases**
 - For e.g., under-represented groups
 - Biased (unfair) performance
 - Fair PCA might be **unstable**
- **Fair VNNs (F-VNNs)**
 - *Fairness:* parity in performance across groups within data



[Cavallo et al., 2025]

Limitations of VNNs - 2

- Datasets may contain harmful **biases**
 - For e.g., under-represented groups
 - Biased (unfair) performance
 - Fair PCA might be **unstable**
- **Fair VNNs (F-VNNs)**



- *Fairness*: parity in performance across groups within data
- How to make VNNs fair?
- Are Fair VNNs stable?

questions to address

[Cavallo et al., 2025]

Fair covariance estimates

➤ Balanced covariance

For two groups g and h ,


$$\hat{\mathbf{C}}_{\text{bal}} = \alpha \hat{\mathbf{C}} + (1 - \alpha)(\hat{\mathbf{C}}_h - \hat{\mathbf{C}}_g) = \alpha_g \hat{\mathbf{C}}_g + \alpha_h \hat{\mathbf{C}}_h$$

➤ Debiased covariance

$$\hat{\mathbf{C}}_{\text{deb}} = \mathbf{X}^T (\mathbf{I}_m + \beta \mathbf{Z} \mathbf{Z}^T)^{-1} \mathbf{X} / n$$

\mathbf{X} : data matrix

\mathbf{Z} : groups of samples



[Cavallo et al., 2025]

Bias-mitigation penalty

- F-VNNs are trained with a **loss penalty** that encourages fairness

$$\min_{\mathcal{H}} \gamma \mathcal{L}(\mathbf{X}, \mathbf{y}, \Phi) + (1 - \gamma) \mathcal{R}(\mathbf{X}, \mathbf{y}, \mathbf{z}, \Phi)$$

\mathcal{L} : task-specific loss (for e.g., cross-entropy, MAE)

\mathcal{R} : bias penalty (for e.g., performance difference across groups)

γ : balancing term

[Cavallo et al., 2025]

Stability of F-VNNs

- Fair covariance estimates
 - $\hat{\mathbf{C}}_{\text{deb}}$ and $\hat{\mathbf{C}}_{\text{bal}}$ are subject to covariance estimation errors
 - PCA with fair covariance estimates (Fair PCA) may be **unstable**
 - ➡ biased treatment

[Cavallo et al., 2025]

Stability of F-VNNs

- Fair covariance estimates
 - $\hat{\mathbf{C}}_{\text{deb}}$ and $\hat{\mathbf{C}}_{\text{bal}}$ are subject to covariance estimation errors
 - PCA with fair covariance estimates (Fair PCA) may be **unstable**
 - ➡ biased treatment

➤ F-VNNs are stable


- Stability of F-VNNs with **balanced** covariance $\propto \mathcal{O}\left(\frac{1}{n_g^{1/2}}\right) + \mathcal{O}\left(\frac{1}{n_h^{1/2}}\right)$
- Stability of F-VNNs with **debiased** covariance $\propto \mathcal{O}\left(\frac{1}{n^{1/2}}\right)$

[Cavallo et al., 2025]

Numerical results

➤ Stability: synthetic biased data

- Train on unbiased dataset
- During test, replace covariance with unbalanced/fair version
- Compare PCA+SVM with VNNs




[Cavallo et al., 2025]

Numerical results

➤ Stability: synthetic biased data

- Train on unbiased dataset
- During test, replace covariance with unbalanced/fair version
- Compare PCA+SVM with VNNs




F-VNNs are more **stable**
F-VNNs achieve **less bias**
F-VNNs **outperform** PCA

[Cavallo et al., 2025]

Numerical results

➤ Real world datasets

Dataset	Description	Task	Sensitive attribute
Parkinson (left)	Medical records of patients	Regression for Parkinson's level	Sex of patient
LSAC (center)	Law school students' features	Regression for GPA	Race of students
German credit (right)	Features of individuals applying for credit	Classification (good or bad)	Sex of individual




F-VNNs achieve better fairness and performance than PCA

[Cavallo et al., 2025]

Limitations-3


- VNN models discussed so far operate on *static* data
 - Real world applications have **dynamic** data
 - Non-trivial modifications needed to handle temporal, non-stationary data
 - Online estimates introduce additional source of errors



Limitations-3

- VNN models discussed so far operate on *static* data
 - Real world applications have **dynamic** data
 - Non-trivial modifications needed to handle temporal, non-stationary data
 - Online estimates introduce additional source of errors

- **Spatio-temporal VNNs (STVNNs)**
VNNs for **spatio-temporal** datasets



Spatiotemporal VNNs

➤ Model design

- Online covariance matrix estimate

$$\hat{\mathbf{C}}_{t+1} = \zeta_t \hat{\mathbf{C}}_t + \beta_t (\mathbf{x}_{t+1})(\mathbf{x}_{t+1})^T$$

[Cavallo et al., 2024]

Spatiotemporal VNNs

➤ Model design


- Online covariance matrix estimate

$$\hat{\mathbf{C}}_{t+1} = \zeta_t \hat{\mathbf{C}}_t + \beta_t (\mathbf{x}_{t+1})(\mathbf{x}_{t+1})^T$$

- Spatio-temporal coVariance filter

$$\mathbf{z}_t := \mathbf{H}(\hat{\mathbf{C}}_t, \mathbf{h}_t, \mathbf{x}_{T:t}) = \sum_{t'=0}^{T-1} \sum_{k=0}^K h_{kt'} \hat{\mathbf{C}}_t^k \mathbf{x}_{t-t'}$$

Spatial and temporal convolution



[Cavallo et al., 2024]

Spatiotemporal VNNs


➤ STVNN

- Sequences of spatio-temporal covariance filters followed by non-linearity

$$\mathbf{z}_t^l = \sigma \left(\mathbf{H}^l(\hat{\mathbf{C}}_t, \mathbf{h}_t, \mathbf{z}_{T:t}^{l-1}) \right)$$

- Online parameter updates

$$\mathbf{h}_{t+1} = \mathbf{h}_t - \eta \nabla_t \mathcal{L}(\Phi(\mathbf{x}_{T:t}, \hat{\mathbf{C}}_t; \mathbf{h}_t))$$



[Cavallo et al., 2024]

Spatiotemporal VNNs

➤ STVNN

- Sequences of spatio-temporal covariance filters followed by non-linearity


$$\mathbf{z}_t^l = \sigma \left(\mathbf{H}^l(\hat{\mathbf{C}}_t, \mathbf{h}_t, \mathbf{z}_{T:t}^{l-1}) \right)$$

- Online parameter updates

$$\mathbf{h}_{t+1} = \mathbf{h}_t - \eta \nabla_t \mathcal{L}(\Phi(\mathbf{x}_{T:t}, \hat{\mathbf{C}}_t; \mathbf{h}_t))$$

- STVNNs are **stable**


$$\text{Stability bound} \propto \mathcal{O} \left(\frac{1}{\sqrt{n}} \right)$$




[Cavallo et al., 2024]

Numerical results


- **Time series forecasting task** (weather data and currency exchange rates)
 - Train with one covariance, test with another estimated from fewer samples



NOAA



Molene




Exchange rate

STVNNs are more stable than temporal-PCA (TPCA)
Higher T (temporal window size), lower stability

Numerical results

- **Time series forecasting task (brain imaging data)**


- **Data:** HCP Young-adult dataset
- BOLD data at spatial scales of 114, 314, 414, 614 (Schaefer's)
- Train model on 314 resolution
 - Test on 114, 414, 614 resolutions



Numerical results

- **Time series forecasting task (brain imaging data)**

- **Data:** HCP Young-adult dataset
- BOLD data at spatial scales of 114, 314, 414, 614 (Schaefer's)
- Train model on 314 resolution
 - Test on 114, 414, 614 resolutions



STVNN demonstrates **transferability** across **multi-scale spatio-temporal** datasets


Conclusions and Future Directions

Covariance filters


- A covariance filter is a **polynomial in the covariance matrix $\hat{\mathbf{C}}$**

$$\mathbf{H}(\hat{\mathbf{C}}) = \sum_{k=0}^K h_k \hat{\mathbf{C}}^k \mathbf{x}$$

- We train the filter coefficients h_k to accomplish some task




CoVariance Neural Networks (VNNs)



- A VNN is a composition of layers
- Each of which is a composition of
 - ... a **covariance filter**
 - ... with a **pointwise nonlinearity**
- $\Phi(\mathbf{x}; \hat{\mathbf{C}}, \mathcal{H})$ represents VNN output
- \mathcal{H} is the set of trainable filter taps

Covariance Filters are Implicitly Equivalent to PCA




Covariance Filters are Implicitly Equivalent to PCA

- The difference is that covariance filters (and VNNs) **do not require eigenvectors**

Stability: Leading to more stable signal processing

Transferability: And the possibility of transferring trained filters across scales



Stability of coVariance filters and Neural Networks

- Outputs of coVariance filters NNs on **true** and **estimated** covariances are close

$$\left\| \mathbf{H}(\hat{\mathbf{C}}) - \mathbf{H}(\mathbf{C}) \right\| = \mathcal{O}\left(\frac{1}{n^{1/2-\varepsilon}}\right) = \alpha_n \quad \left\| \Phi(\mathbf{x}, \hat{\mathbf{C}}; \mathcal{H}) - \Phi(\mathbf{x}, \mathbf{C}; \mathcal{H}) \right\| \leq LF^{L-1} \alpha_n$$

- Provided that the filters (at each layer) have Lipschitz frequency responses

$$|h(\lambda_i) - h(\lambda_j)| \leq Q \frac{|\lambda_i - \lambda_j|}{k_i}$$

- This requirement limits discriminability but it is a necessary limit


[Sihag et al., 2022]

PCA responds catastrophically to eigenspace estimation

- Difference between **true** and **estimated** eigenvectors can be arbitrarily different

$$\|\hat{\mathbf{V}}_{\mathbf{x}} - \mathbf{V}_{\mathbf{x}}\| = \mathcal{O}\left(\frac{1}{n^{1/2} \min_{i \neq j} |\lambda_i - \lambda_j|}\right)$$

- Filters process **similar eigenvalues** with **similar coefficients**





PCA responds catastrophically to eigenspace estimation

- Difference between **true** and **estimated** eigenvectors can be arbitrarily different

$$\|\hat{\mathbf{V}}_{\mathbf{x}} - \mathbf{V}_{\mathbf{x}}\| = \mathcal{O}\left(\frac{1}{n^{1/2} \min_{i \neq j} |\lambda_i - \lambda_j|}\right)$$

- Eigenvectors with **similar eigenvalues** can be processed with **dissimilar coefficients**




PCA responds catastrophically to eigenspace estimation

- Difference between **true** and **estimated** eigenvectors can be arbitrarily different


$$\|\hat{\mathbf{V}}_{\mathbf{x}} - \mathbf{V}_{\mathbf{x}}\| = \mathcal{O}\left(\frac{1}{n^{1/2} \min_{i \neq j} |\lambda_i - \lambda_j|}\right)$$

- Filters process **similar eigenvalues** with **similar coefficients**




coVariance Filters and VNNs are Scale-Free Models

- Filters and VNNs are defined by coefficients that we can **transfer across scales**
 - Train at small scale and transfer to large scale
 - Train jointly across a heterogeneous range of scales




VNNs are provably transferable



Transferability bound

$$\|y_{m_1} - y_{m_2}\| \propto \mathcal{O}\left(\frac{1}{m_1^{3\zeta/2-1}} + \frac{1}{m_2^{3\zeta/2-1}}\right), \text{ for } \zeta \in (2/3, 1]$$


VNNs are provably transferable to limit models



Transferability bound* [Sihag et al., 2024]

$$\|y_{m_1} - y\| \propto \mathcal{O}\left(\frac{1}{m_1^{3\zeta/2-1}}\right), \text{ for } \zeta \in (2/3, 1]$$

*Assumption: data is a discretization of a common continuous model




VNNs are well suited for neuroimaging data analysis


- Properties of VNNs make them appealing for neuroimaging data analysis
- **Connections with PCA** → transparent outcomes by leveraging spectrum of covariance matrix
- **Stability** → reproducible outcomes in limited data settings
- **Transferability** → enhanced generalizability and robustness to choice of brain atlases

Brain age gap prediction is a transfer learning problem

- Train ML model to predict age on a **large dataset (healthy population)**



- Apply the **pre-trained** ML model on a **target dataset (neurodegeneration)**




- Brain age gap is the **residual** of the model


A principled approach to brain age gap prediction

- **Focus on residuals** of the ML model, not prediction performance
- **Qualitative evaluation** during pre-training
 - what does the model learn during **pre-training** on **healthy population**?
- **Interpretability/explainability:**
 - what's driving elevated brain age gap (residuals) in **neurodegeneration**?
- **Generalizability** to diverse target populations


VNNs provide an anatomically interpretable and explainable brain age gap



VNNs provide an anatomically interpretable and explainable brain age gap



VNNs provide an anatomically interpretable and explainable brain age gap



Concluding Remarks

➤ Emerging areas

- **Sparse VNNs:** sparsifying covariance matrix [Cavallo et al., 2024]
- **Spatiotemporal VNNs:** temporal datasets [Cavallo et al., 2024]
- **Fair VNNs:** unbiased outcomes with VNNs [Cavallo et al., 2025]
- **Optimality of covariance matrices:** suitability of covariance to learning task
[Khalafi et al., 2024]
- **Application to brain age gap prediction** [Sihag et al., 2024; 2025]

Future directions

- Learning with **cross-covariance** graphs
 - Links with partial least squares/ canonical correlation analysis
- Expand **interpretability/explainability** of VNNs
 - How are eigenvectors exploited in STVNNs on dynamic datasets?
- Building **interpretable biomarkers**
 - Using other modalities (for e.g., fMRI)

References

- Sihag, Saurabh, Mateos, Gonzalo, C. McMillan, and Ribeiro, Alejandro, "coVariance neural networks," in Proc. Conference on Neural Information Processing Systems, Nov. 2022.
- Saurabh Sihag, Gonzalo Mateos, C. McMillan, and Alejandro Ribeiro, "Explainable brain age prediction using covariance neural networks," in Proc. Conference on Neural Information Processing Systems, 2023.
- Saurabh Sihag, Gonzalo Mateos, and Alejandro Ribeiro, "Disentangling neurodegeneration with brain age gap prediction models: A graph signal processing perspective," in IEEE Signal Processing Magazine, 2025 (to appear).
- Sihag, Saurabh, Mateos, Gonzalo, C. McMillan, and Ribeiro, Alejandro, "Transferability of covariance neural networks," in IEEE Journal of Selected Topics in Signal Processing, pp. 1–16, 2024.
- S. Khalafi, Saurabh Sihag, and Alejandro Ribeiro, "Neural tangent kernels motivate cross-covariance graphs in neural networks," in International Conference on Machine Learning, 2024.
- Sihag, Saurabh, Mateos, Gonzalo, and Ribeiro, Alejandro, "Explainable brain age gap prediction in neurodegenerative conditions using covariance neural networks," IEEE International Symposium on Biomedical Imaging, 2025.
- A. Cavallo, Z. Gao, and Elvin Isufi, "Sparse covariance neural networks," arXiv:2410.01669, vol. cs.LG, 2024.
- Cavallo, Andrea, et al. "Fair covariance neural networks." IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP). IEEE, 2025.
- A. Cavallo, M. Sabbaqi, and Isufi, Elvin, "Spatiotemporal covariance neural networks," in Joint European Conference on Machine Learning and Knowledge Discovery in Databases, pp. 18–34, Springer, 2024.

Slides available at

